

# Linux Firewall Exploration Lab

Xinyi Li

March 30, 2020

Instruction: [https://seedsecuritylabs.org/Labs\\_16.04/PDF/Firewall.pdf](https://seedsecuritylabs.org/Labs_16.04/PDF/Firewall.pdf)

Set up 2 VMs:

- A: 10.0.2.15
- B: 10.0.2.4.

## Task 1

All commands run on the VM 10.0.2.15, and modify Line 11 in its /etc/default/ufw with root privilege (e.g. sudo vi/sudo gedit/sudo nano+filename or whatever method you like) as:

```
1 DEFAULT_INPUT_POLICY="ACCEPT"
```

Or simply use the command:

```
1 sudo ufw default allow incoming
```

Ref to the manual and tutorial of ufw

At first, enable it:

```
1 sudo ufw enable
```

After configuration, reset it to the installed status and remove added rules:

```
1 sudo ufw reset
```

### Prevent A from doing telnet to Machine B

```
1 sudo ufw deny out from 10.0.2.15 to 10.0.2.4 port 23
```

### Prevent B from doing telnet to Machine A

```
1 sudo ufw deny in from 10.0.2.4 to 10.0.2.15 port 23
```

## Prevent A from visiting an external web site

Use ping or traceroute to get one of the host IP address and block the corresponding HTTP/HTTPS connections:

```
1 sudo ufw deny out from 10.0.2.15 to host_ip port 80
2 sudo ufw deny out from 10.0.2.15 to host_ip port 443
```

**Note:** As this answer explains, It is impossible to stop the accessing to a domain.

## Task 2

From the programming manual of `netfilter` module, I can implement a simplified firewall program as `packet_filter.c`

For each rule, a callback function is defined to filter packets meeting some specified conditions.

For example, the function for task 1.1 can be defined as `telnetFilter_1()`

```
1 unsigned int telnetFilter_1(void *priv, struct sk_buff *skb,
2                             const struct nf_hook_state *state)
3 // rule for task 1.1: Prevent A from doing `telnet` to Machine B
4 {
5     struct iphdr *iph;
6     struct tcphdr *tcpiph;
7
8     iph = ip_hdr(skb);
9     tcpiph = (void *)iph + iph->ihl * 4;
10
11    if (iph->protocol == IPPROTO_TCP && tcpiph->dest == htons(23)
12        && eq_daddr(iph, "10.0.2.4") && eq_saddr(iph,
13                  "10.0.2.15"))
14    {
15        printk(KERN_INFO "Dropping telnet from %pI4 packet to
16              %pI4\n", &iph->saddr, &iph->daddr);
17        return NF_DROP;
18    }
19    else
20    {
21        return NF_ACCEPT;
22    }
23 }
```

Similarly, the `if` statement can be replaced by other rules to construct more filters.

The `if` condition in `telnetFilter_2()` for task 1.2:

```

1 if (iph->protocol == IPPROTO_TCP && tcph->dest == htons(23) &&
    eq_daddr(iph, "10.0.2.15") && eq_saddr(iph, "10.0.2.4"))

```

Assume that we intend to block machine A from opening the website: <http://notebook.xyli.me/>. Before constructing the rule, we utilize Wireshark to get its 2 host IP address: 104.18.21.226 and 103.235.46.191.

Based on the knowledge, the `if` condition in ‘block\_xyli\_me’ for task 1.3 should be:

```

1 if ((tcph->dest == htons(80) || tcph->dest == htons(443))
2 && (eq_daddr(iph, "104.18.21.226") ||
     eq_daddr(iph, "103.235.46.191"))
3 && eq_saddr(iph, "10.0.2.15"))

```

To make it scalable for at least 5 filter rules, a `nf_hook_ops` array should be declared with a large size and `regist_num` is maintained to track the actual amount of filters used in the firewall:

```

1 #define MAX_RULE_NUM 10
2
3 static struct nf_hook_ops FilterHookRule[MAX_RULE_NUM];
4 static int regist_num = 0;

```

Regist those hooks with functions above in `functionsetUpFilter()`:

```

1 int setUpFilter(void)
2 {
3     int i;
4     printk(KERN_INFO "Registering filters.\n");
5     FilterHookRule[0] = (struct nf_hook_ops){.hook =
6         telnetFilter_1, .hooknum = NF_INET_LOCAL_OUT, .pf =
7         PF_INET, .priority = NF_IP_PRI_FIRST};
8     FilterHookRule[1] = (struct nf_hook_ops){.hook =
9         telnetFilter_2, .hooknum = NF_INET_LOCAL_IN, .pf =
10        PF_INET, .priority = NF_IP_PRI_FIRST};
11     FilterHookRule[2] = (struct nf_hook_ops){.hook =
12        block_xyli_me, .hooknum = NF_INET_LOCAL_OUT, .pf =
13        PF_INET, .priority = NF_IP_PRI_FIRST};
14
15     // set the amount of filter rules
16     regist_num = 3;
17
18     for (i = 0; i < regist_num; i++)
19         nf_register_hook(&FilterHookRule[i]);
20     return 0;
21 }

```

When extending the module with more rules, just focus to fill out `hooknum` and `hook`(i.e. function definition) fields like this.

**Note:** `hooknum` field, namely hook type, is not consistent with the identifiers in the book. Please read their actual definition (alias `enum`) in `linux/netfilter.h`.

Define an unreg function as well and associate those functions in the module.

```
1 void removeFilter(void)
2 {
3     int i;
4     printk(KERN_INFO "Filters are being removed.\n");
5     //unregist hooks one by one
6     for (i = 0; i < regist_num; i++)
7         nf_unregister_hook(&FilterHookRule[i]);
8     regist_num = 0;
9 }
10
11 module_init(setUpFilter);
12 module_exit(removeFilter);
13
14 MODULE_LICENSE("GPL");
```

Finally, write a `makefile` and type `make` to compile:

```
1 obj-m += packet_filter.o
2 all:
3     make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
4
5 clean:
6     make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

Install the output module `packet_filter.ko` into the kernel:

```
1 sudo insmod packet_filter.ko
```

Now, you can observe the firewall works, and log can be viewed with command `dmseg`:

```
t to 103.235.46.191
[10221.204003] Dropping http/https from 10.0.2.15 packe
t to 103.235.46.191
[10221.459903] Dropping http/https from 10.0.2.15 packe
t to 103.235.46.191
[10236.300782] Dropping http/https from 10.0.2.15 packe
t to 103.235.46.191
[10236.300789] Dropping http/https from 10.0.2.15 packe
t to 103.235.46.191
[10236.300791] Dropping http/https from 10.0.2.15 packe
t to 103.235.46.191
[10236.300793] Dropping http/https from 10.0.2.15 packe
t to 103.235.46.191
[10237.323890] Dropping http/https from 10.0.2.15 packe
t to 103.235.46.191
[10237.579689] Dropping http/https from 10.0.2.15 packe
t to 103.235.46.191
[10271.098951] Dropping http/https from 10.0.2.15 packe
t to 103.235.46.191
[10271.098954] Dropping http/https from 10.0.2.15 packe
[03/28/20]seed@VM:~/.../filter$
```

Don't forget to uninstall the firewall after the lab:

```
1 sudo rmmod packet_filter
```

## Task 3

**Block all the outgoing traffic to external telnet servers.**

Set up another VM C: 10.0.2.5.

On machine A:

```
1 sudo ufw deny out from 10.0.2.15 to any port 23
2 sudo ufw enable
```

**Block all the outgoing traffic to www.facebook.com**

~~It's difficult because the domain has many dynamic IP addresses as hosts. For example, ping it at different times:~~

```
1 $ping www.facebook.com
```

```

2 PING star-mini.c10r.facebook.com (157.240.2.35) 56(84) bytes of
   data.
3 64 bytes from edge-star-mini-shv-01-ort2.facebook.com
   (157.240.2.35): icmp_seq=1 ttl=52 time=18.7 ms

1 $ ping www.facebook.com
2 PING star-mini.c10r.facebook.com (157.240.18.35) 56(84) bytes of
   data.
3 64 bytes from edge-star-mini-shv-02-ort2.facebook.com
   (157.240.18.35): icmp_seq=1 ttl=52 time=19.3 ms

```

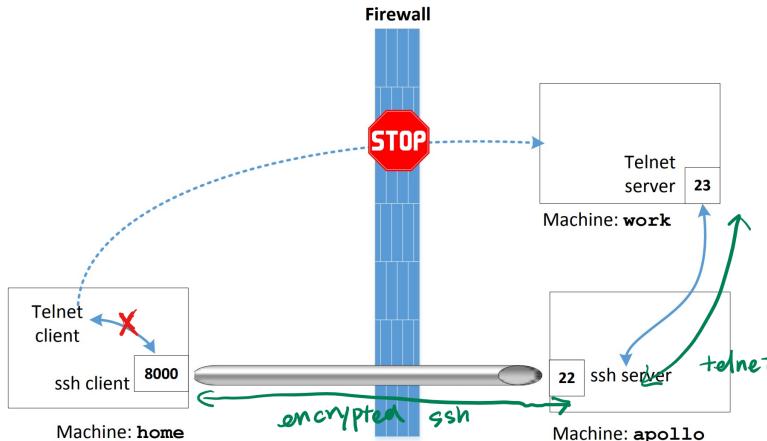
To ensure the blocking, you can use the subnet instead of a specified IP address:

```
1 sudo ufw deny out from 10.0.2.15 to 157.240.2.35/16
```

### Task 3.a: Telnet to Machine B through the firewall

**Note:** The solution given by the instruction may not work. Perhaps because `telnet` usually fails on ports other than a few certain ports (see <https://aplawrence.com/DirkHart/telnet2023.html> and <https://superuser.com/questions/790782/using-telnet-port-number>)

Here is a simple alternative method:

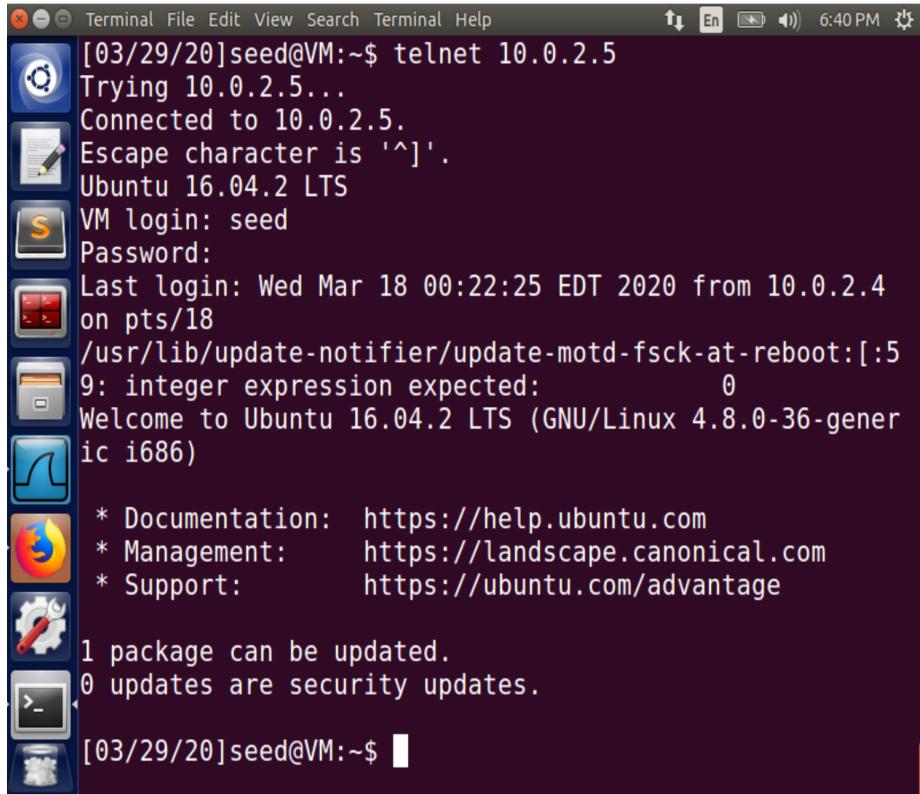


Don't do any extra port mappings. use the most naive way to `ssh` machine B

```
1 ssh seed@10.0.2.4
```

type the password and directly `telnet` machine C

```
1 telnet 10.0.2.5
```



A screenshot of a Linux desktop environment, likely Ubuntu 16.04 LTS, showing a terminal window. The terminal window title is "Terminal". The terminal content shows a telnet session to 10.0.2.5:

```
[03/29/20]seed@VM:~$ telnet 10.0.2.5
Trying 10.0.2.5...
Connected to 10.0.2.5.
Escape character is '^].
Ubuntu 16.04.2 LTS
VM login: seed
Password:
Last login: Wed Mar 18 00:22:25 EDT 2020 from 10.0.2.4
on pts/18
/usr/lib/update-notifier/update-motd-fsck-at-reboot:[ :5
9: integer expression expected:          0
Welcome to Ubuntu 16.04.2 LTS (GNU/Linux 4.8.0-36-generic i686)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

1 package can be updated.
0 updates are security updates.

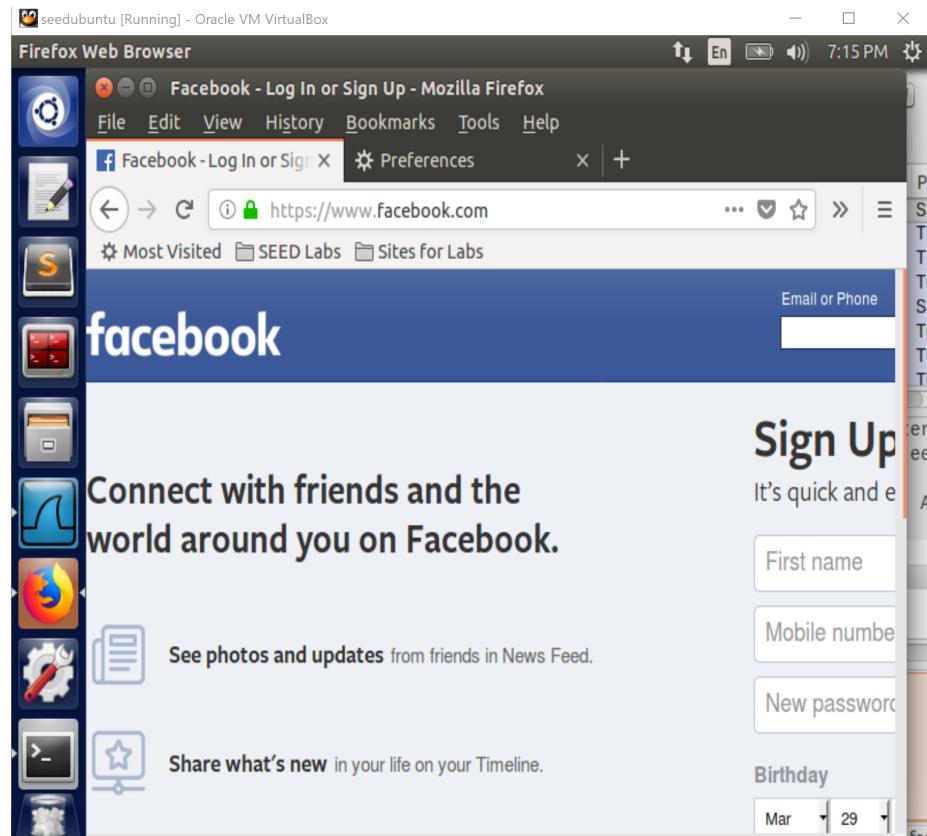
[03/29/20]seed@VM:~$
```

### Task 3.b: Connect to Facebook using SSH Tunnel

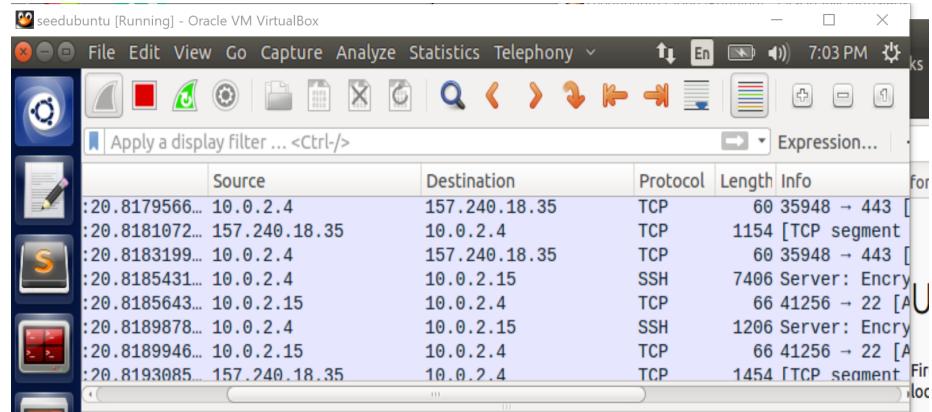
```
1 ssh -D 500 -C seed@10.0.2.4
```

Then configure on Firefox, change the network proxy as manual proxy mode with SOCKS host: 127.0.0.1:9000

1. Now, I can open [www.facebook.com](http://www.facebook.com)



2. When the `ssh` connection is broken, clean the history data in the browser, the website `www.facebook.com` cannot be accessed any more.
3. re-connect to machine B with `ssh` command above, the website `www.facebook.com` can be opened now.
4. And using Wireshark, it shows that the actual HTTP packet is communicated between the external host of `facebook` and machine B, then our local machine A communicate with machine B by `SSH` to request the data from `facebook`:



## Task 4

Assume that VM A is the internal machine which blocks `ssh` connections requested from VM B.

Therefore, on VM A, run

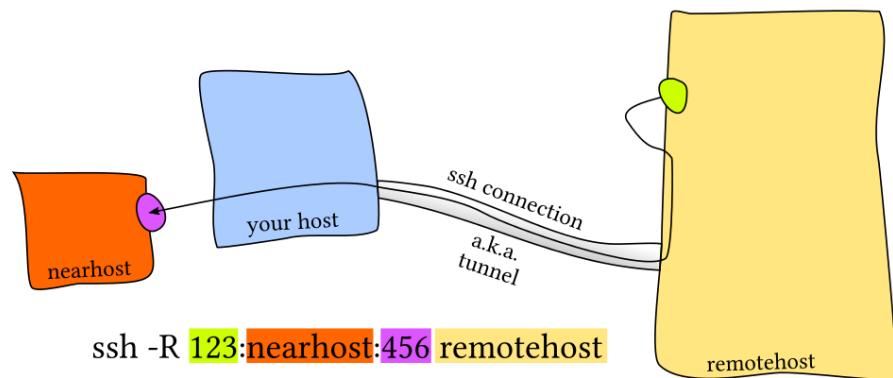
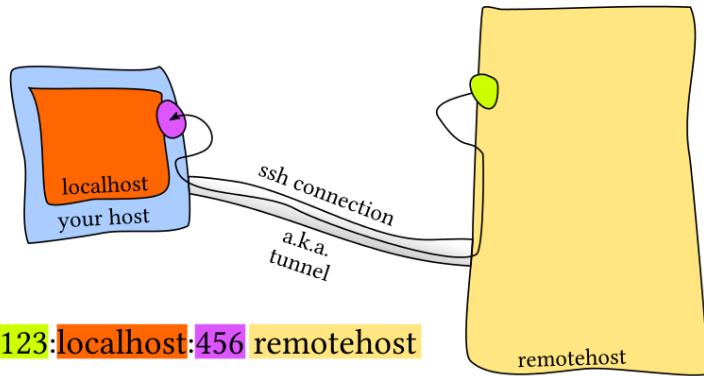
```
1 sudo ufw deny in from 10.0.2.4 to 10.0.2.15 port 22
2 sudo ufw deny in from 10.0.2.4 to 10.0.2.15 port 80
```

to emulate the set-up.

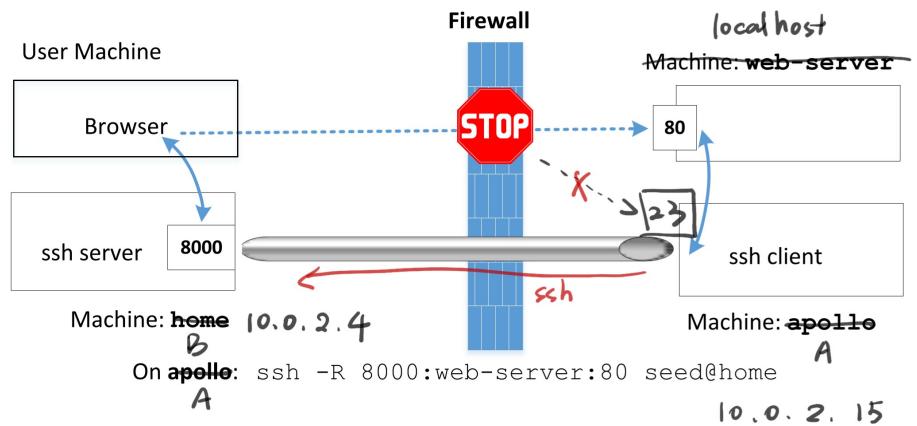
Then we can create a reverse `ssh` tunnel on machine A:

```
1 ssh -R 8000:localhost:80 seed@10.0.2.4
```

Reference to How does reverse SSH tunneling work?



Keep the `ssh` connection alive and leave machine A alone.



We can access the web server hosted by machine A by `localhost:8000` on machine B now:

